



CS326 – Systems Security


Lecture 10

Introduction to Software Security

Elias Athanasopoulos
athanasopoulos.elias@ucy.ac.cy

Recall

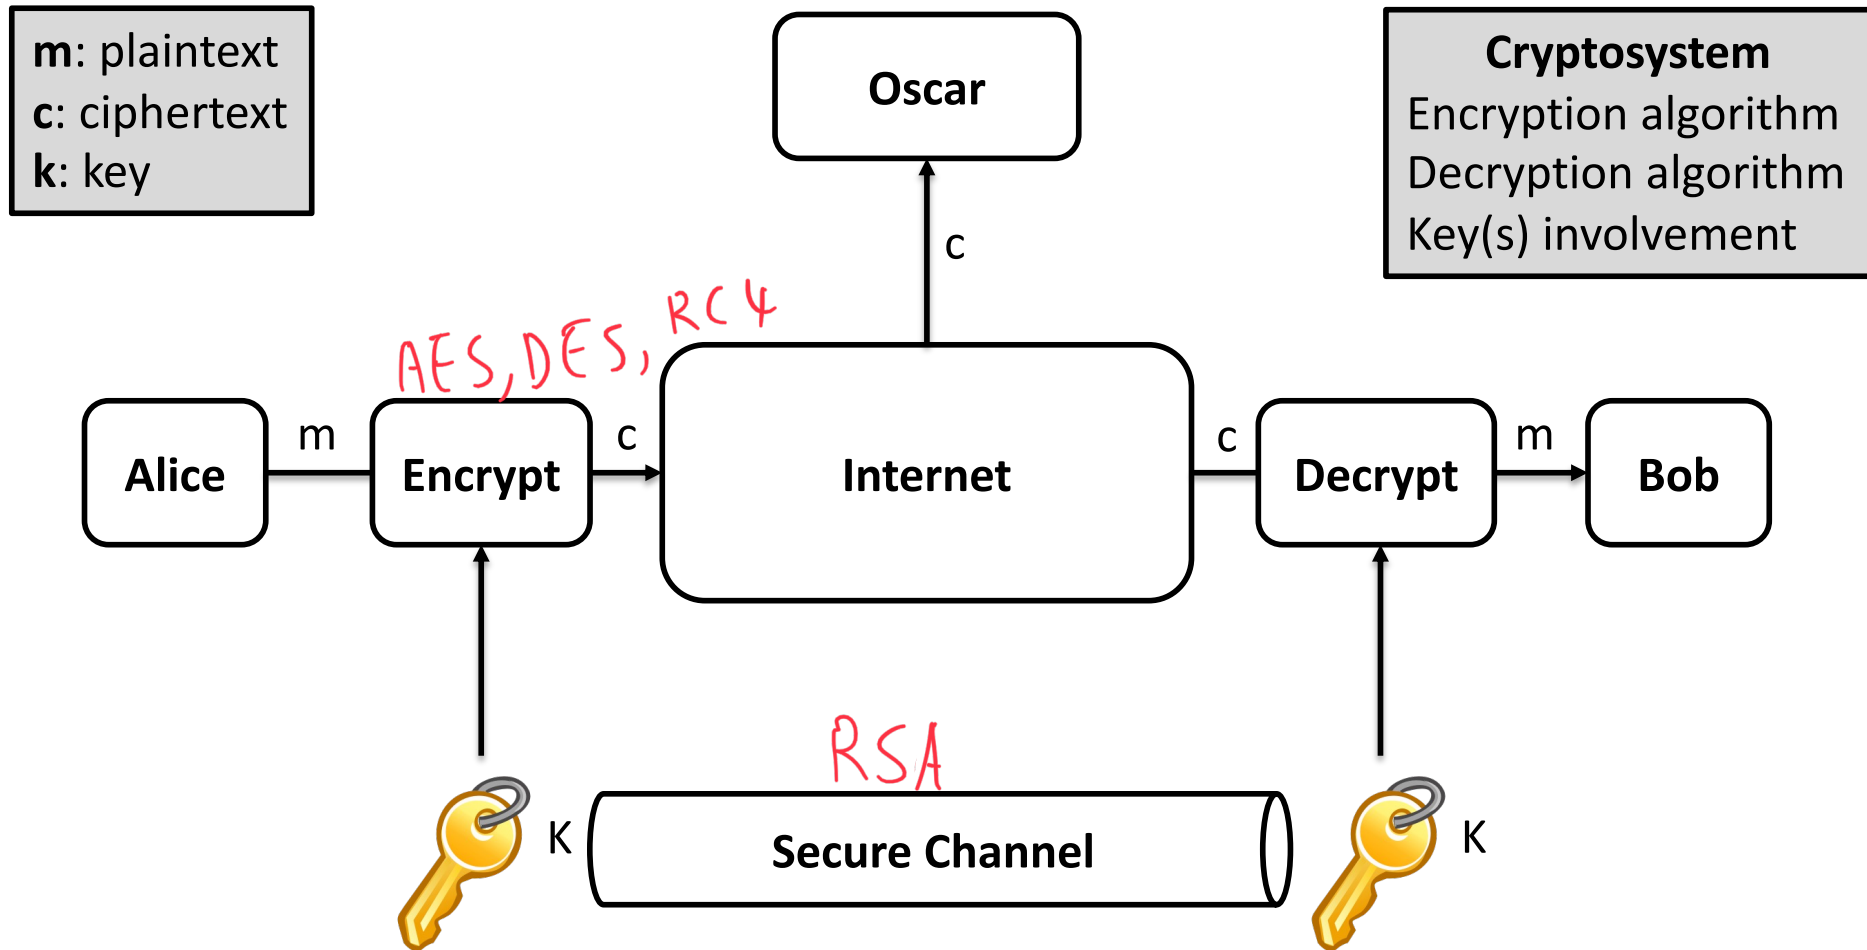
Basic Problem



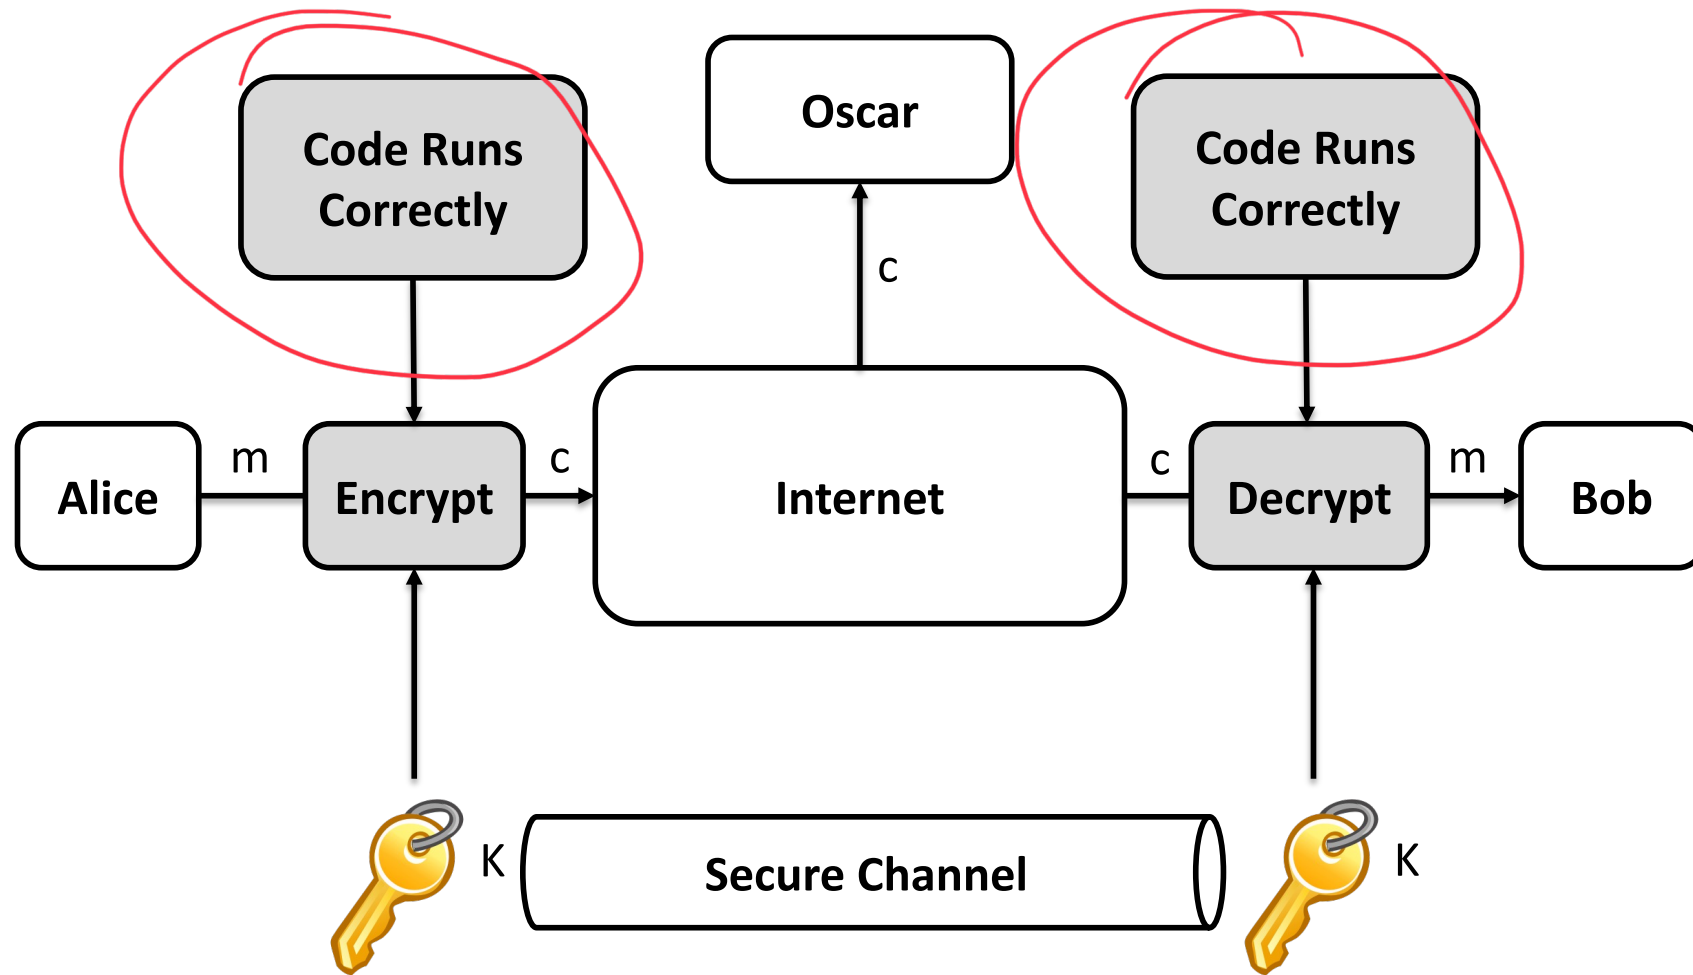
Oscar can see the message (confidentiality)
Oscar can modify the message (integrity)

Recall

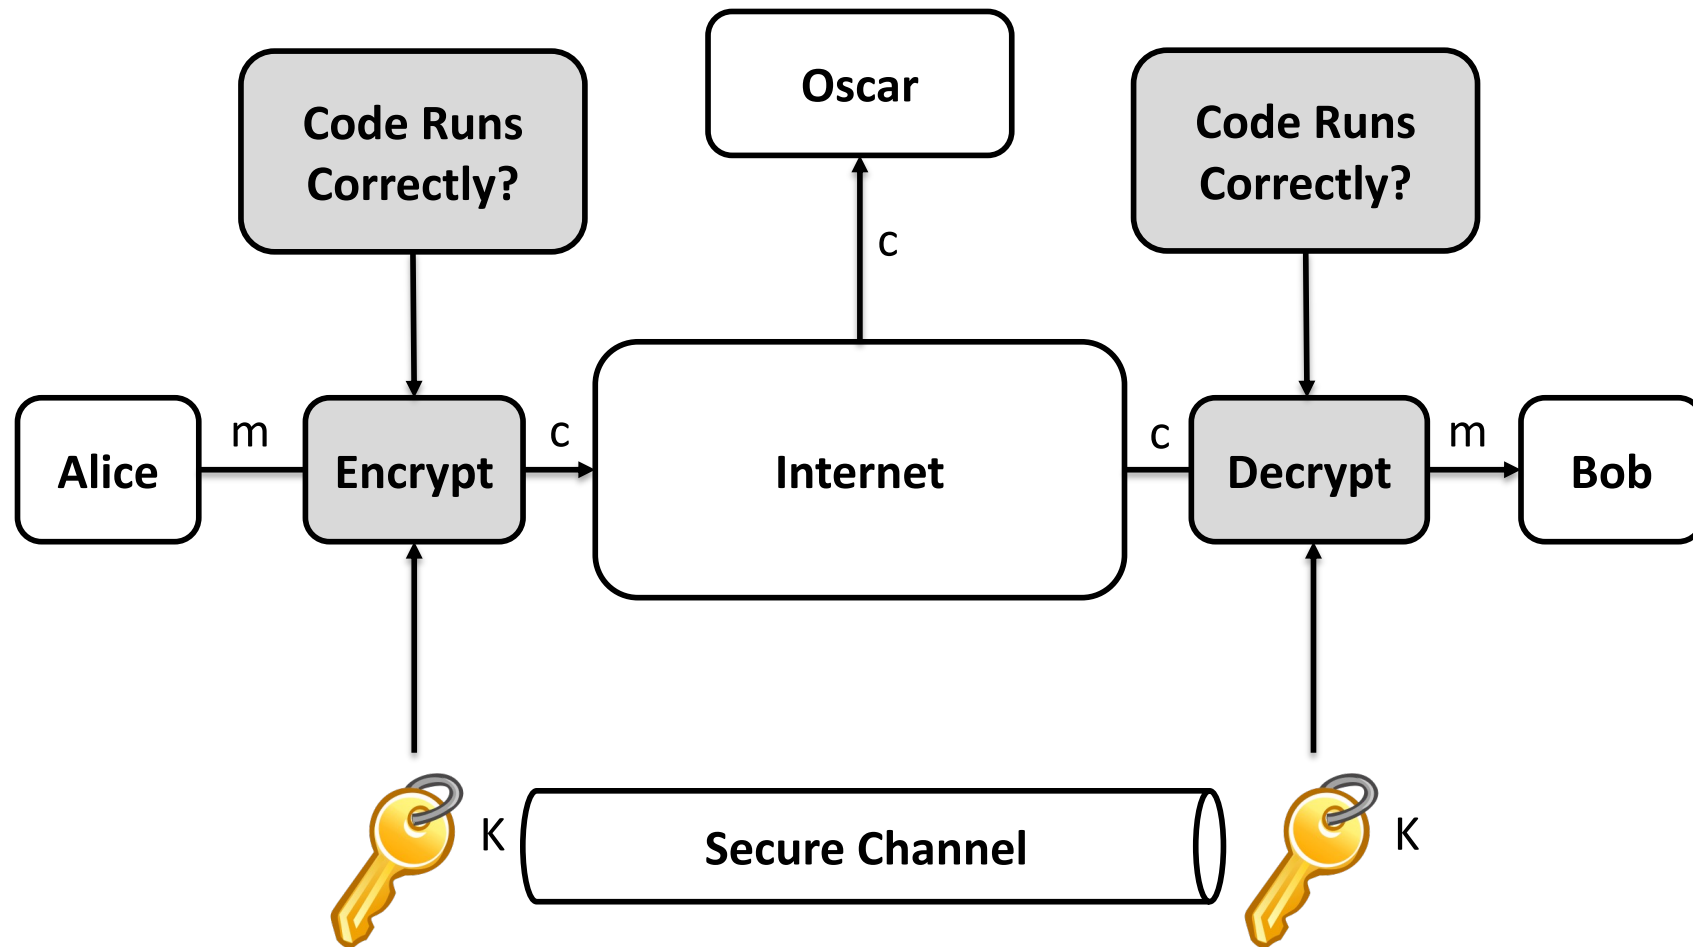
Cryptographic Approach



Implicit Assumption



Basic Problem



Software Roadmap



- Software properties
- Software bugs
- Attacks and exploitation
- Defenses

Software



- Programs are written in a high-level language
 - C, C++, Java, C#, Ruby, Python
- They are compiled for execution
 - Machine code (unmanaged and unsafe code)
 - Virtual-machine code, such as JVM (managed and safe code)
- Different architectures exhibit different properties in executing software
 - Some generic concepts apply to all

Safe vs Unsafe Systems




- Safe programming systems
 - Execute managed code in a virtual machine (e.g., Java)
 - Perform static analysis and compile time (reject unsafe code) and use run-time checks (e.g., Rust)
 - Restricted memory access
- Unsafe programming systems
 - Unrestricted memory access (e.g., C/C++)
 - Better performance
 - Low-level accessing (e.g., drivers)
 - Legacy code

Unsafe vs Unsafe Programming Systems



Unsafe (e.g., C/C++)

```
a[i] = j;
```




Machine Code

- `a` is just a starting address
- No way to check if `i` is in the bounds of `a`

Safe (e.g., Java)

```
a[i] = j;
```



Virtual Machine Code

- `a` is a well-defined entity
- VM knows `a`'s type, bounds, reference counts, etc.

Software Composition



- Data
 - Static pre-defined and allocated at compile-time
 - Dynamically allocated and short-lived (usually at the stack of the program)
 - Dynamically allocated and long-lived (usually at the heap of the program)
- Code
 - Known at compile-time (ahead of time)
 - Generated at run-time (just-in time) using a JIT engine

Compilation, Loading and Execution



- The compiler prepares the program
 - Creates a binary file stored in the disk
 - Binary contains code and data
 - Binary has different sections
 - Binary may have additional dependencies (shared libraries)
- The loader reads the binary and maps it to memory
 - Resolves symbols
 - Loads shared libraries
- The OS creates a structure called *process*
 - A process is a virtual concept of an executing program
 - A binary (stored in the disk) has a different layout from an executing process (mapped in memory)

ld linux.so

Processes



- Programs are executed with the help of the OS
- The OS creates an internal structure (called virtual process) for running the program
- Each virtual process has several things
 - Resources (opened files and sockets, mapped files, etc.)
 - State (page tables, register values, memory contents, etc.)
 - A pointer to the next instruction (program counter)
- Processes execute as they are alone
 - All (virtual) memory is available for a running process
 - The OS can execute several processes, concurrently

Virtual Memory



- A process works using memory
- Memory is given by the OS in a number of virtual pages (a few kilobytes each)
- Physical memory is mapped to virtual pages
 - Every process has its own *map* (page tables)
 - If physical memory runs out, it is swapped to disk
- Not all pages are loaded at once
 - Code pages are loaded when needed (using *pagefaults*)
 - Additional pages for data are allocated upon request
- Accessing a memory address that is not contained in a loaded page causes a pagefault
 - Unless the OS can load the page (i.e., a code page) or handle it somehow, the process will crash

Virtual Memory



- It is practical to have different memory areas
 - Code region has pages having code
 - Data region has pages having data
 - Stack has pages hosting temporary data
 - Heap has pages with dynamically allocated data
- Pages have permissions (rwx)
 - Some enforced by the hardware (MMU)
 - NX-bit → non-executable bit

Example of a C program



```
#include <stdio.h>
```

```
int main(int argc, char *argv[]) {
```

```
    int ary[5] = {1, 2, 3, 4, 5}; → stack
```

```
    fprintf(stderr, "The fifth number of ary is: %d\n", ary[5]);
```

```
    return 1;
```

```
}
```

without stdio.h, stderr(2) will not be found

Out of bounds access!



```
#include <stdio.h>

int main(int argc, char *argv[]) {

    int ary[5] = {1, 2, 3, 4, 5};

    fprintf(stderr, "The fifth number of ary is: %d\n", ary[5]);

    return 1;
}
```


Vocabulary



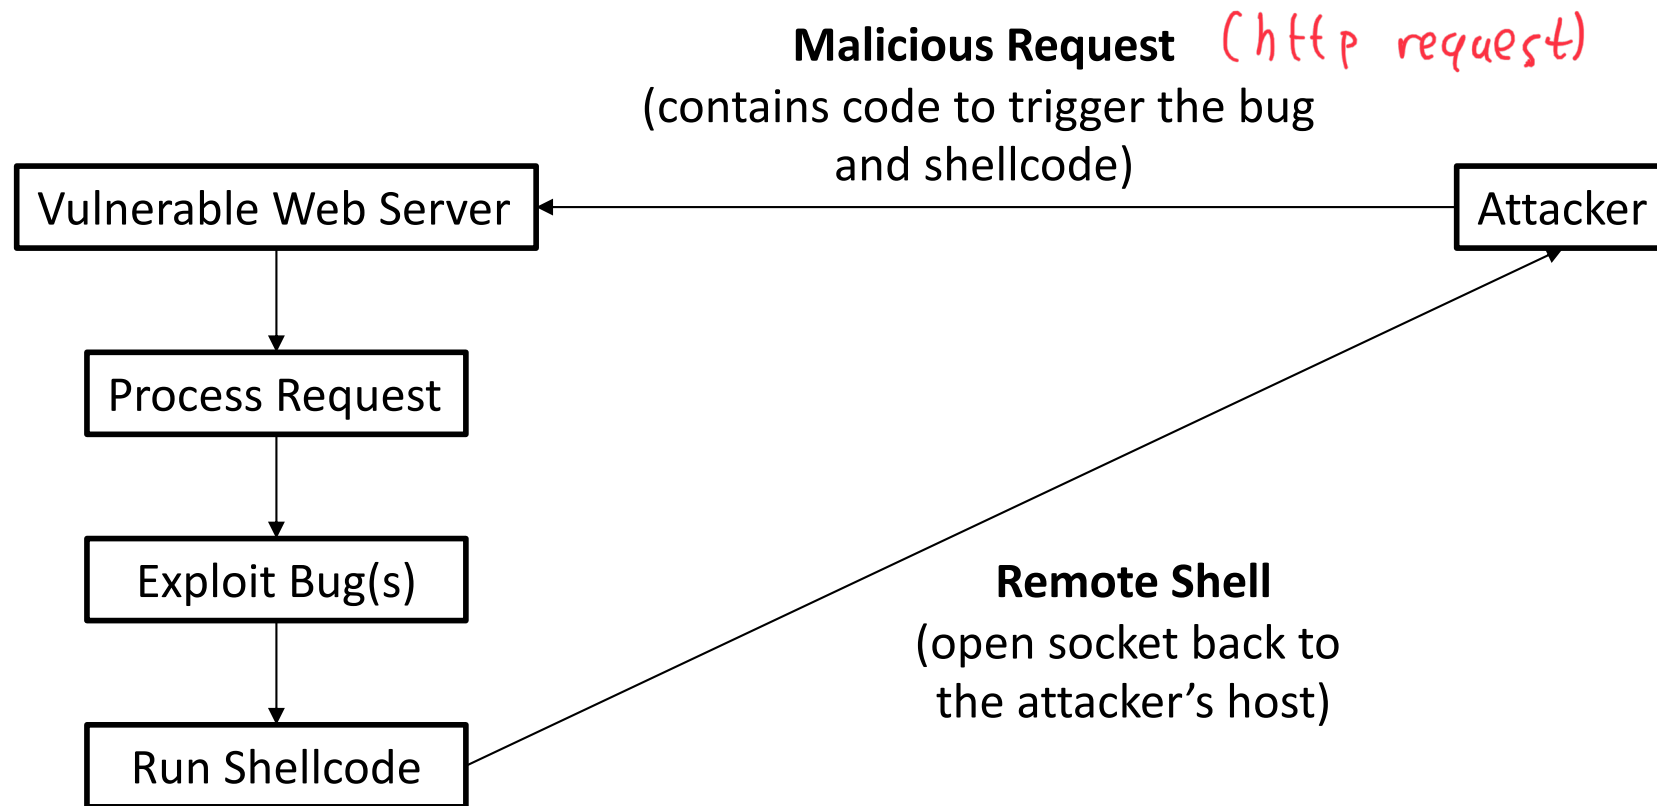
- Vulnerability (bug)
 - A software error (also known as bug) that can potentially allow someone to take advantage of the vulnerable program
- Exploit (process to take advantage of the bug)
 - The process of controlling a program by taking advantage of one or more vulnerabilities
 - Not all vulnerabilities can be exploited

Vocabulary

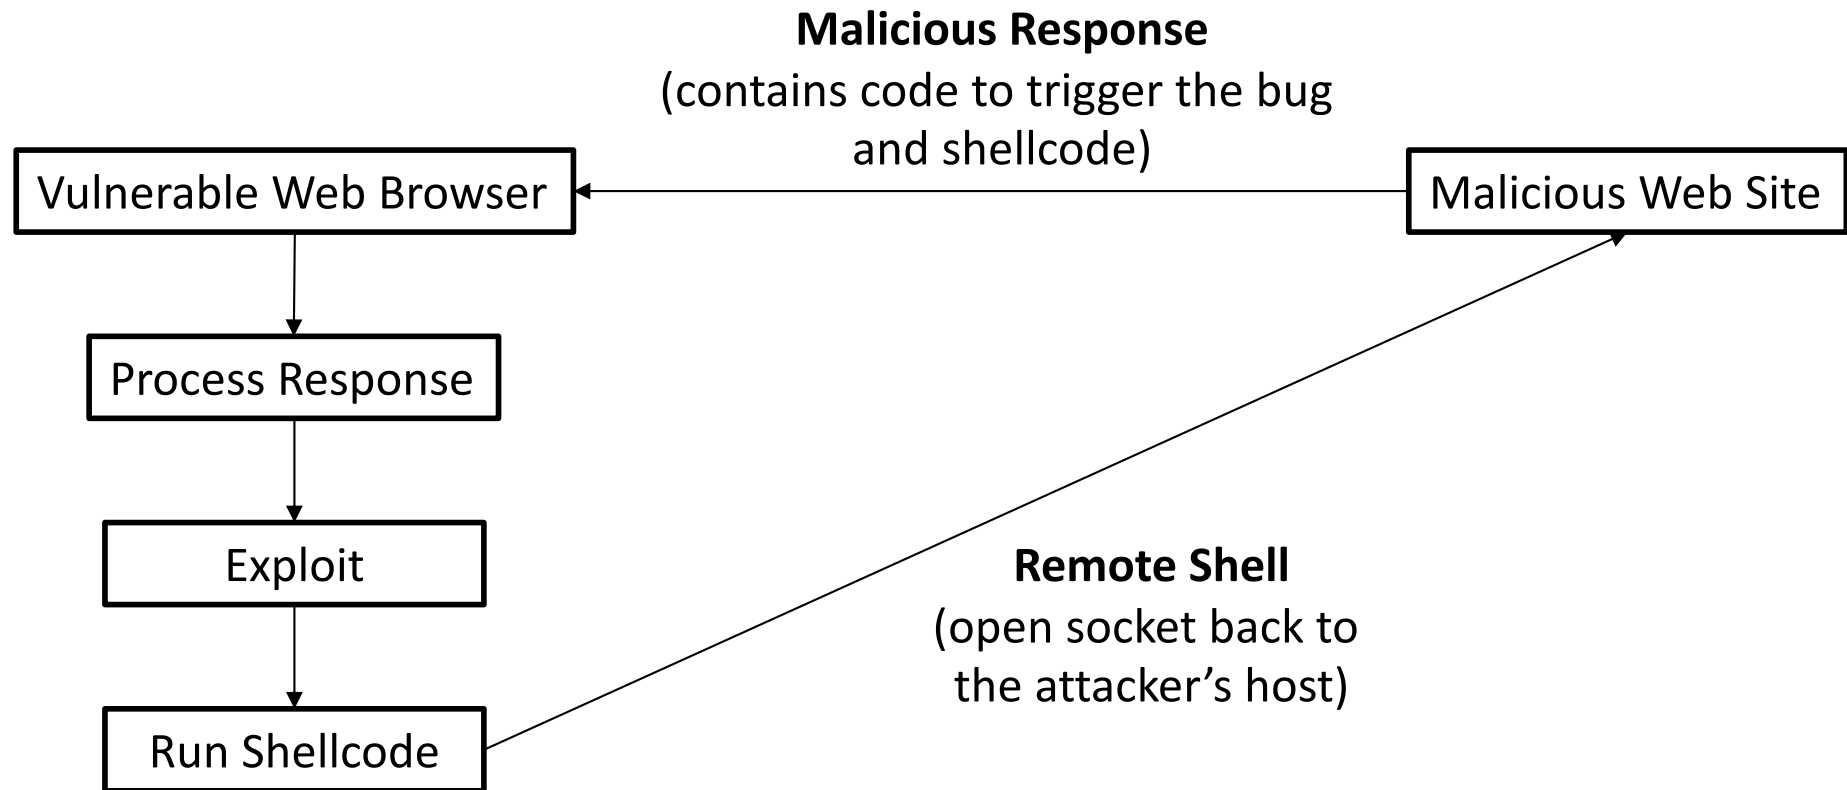


- Arbitrary Code Execution
 - The state of an exploit where an attacker can execute a program of their choice
- Shellcode
 - A machine code that a vulnerable program executes and serves the purposes of the attacker
 - Spawn a shell (can be remote), download malware, create a hidden account, manipulate software, etc.
 - Heavily architecture dependent

High-level Idea



High-level Idea



Software Exploitation



- The victim program has vulnerabilities
 - Can be a program executing in user-space
 - Can be the OS
- Bugs can be triggered using malicious inputs
 - Inputs can be sent over the network (remote attacker)
 - Inputs can be sent locally (local attacker)
- Triggering the bugs can lead to arbitrary code execution
- Arbitrary code execution can run the shellcode